


# HOW TO BUILD A NON-VOLATILE MEMORY DATABASE SYSTEM

---


JOY ARULRAJ  
CARNEGIE MELLON UNIVERSITY



# NON-VOLATILE MEMORY (NVM)



DRAM



NVM



SSD

*Like DRAM, low latency loads and stores*

*Like SSD, persistent writes and high density*

*Why we think NVM is happening  
for real this time?*

# #1: INDUSTRY STANDARDS

- Standard definitions of NVM technologies
  - *Form factors (e.g., JEDEC classification)*
  - *Interface specifications (e.g., NVM Express over Fabrics)*


JUNE 2016



## #2: OPERATING-SYSTEM SUPPORT

- Growing OS support for NVM
  - *Linux 4.8, e.g. NVM Express over Fabrics library*
  - *Windows 10, e.g. Direct access to files on NVM*

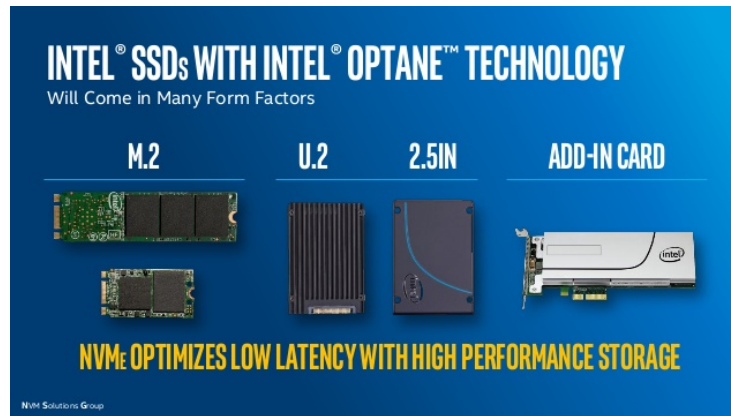
OCTOBER 2016



# #3: PROCESSOR SUPPORT

- Intel's Kaby Lake processor
  - *Support for 3D XPoint NVM technology*
  - *ISA updates for NVM management*


**JANUARY 2017**



*How can we leverage NVM in a DBMS?*

# #1: DISK-ORIENTED DBMSs

- Treat NVM like a faster SSD




*Designed to minimize  
random writes to NVM*

*But, NVM supports  
fast random writes*



## #2: MEMORY-ORIENTED DBMSs

- Treat NVM as extended memory




*Designed to overcome the  
volatility of memory*


*But, writes to NVM  
are persistent*



**DBMS  
OVERVIEW**



**LOGGING &  
RECOVERY**



**DATA  
PLACEMENT**

# NVM-AWARE DBMS OVERVIEW

**EXECUTION  
ENGINE**

**PLAN  
EXECUTOR**

**QUERY  
OPTIMIZER**

**SQL  
EXTENSIONS**

**STORAGE  
MANAGER**

**LOGGING &  
RECOVERY**

**DATA  
PLACEMENT**

**ACCESS  
METHODS**

**ACCESS  
INTERFACES**

**ALLOCATOR  
INTERFACE**

**FILESYSTEM  
INTERFACE**




HOW TO BUILD A NON-VOLATILE MEMORY DBMS  
UNDER SUBMISSION

# #1: ACCESS INTERFACES

---

- Allocator Interface
  - *Provide a durability primitive*
  - *Prevent persistent memory leaks*
- Filesystem Interface
  - *Direct access to files on NVM*
  - *Avoid extra copy in page cache*



## #2: STORAGE MANAGER

---

- Logging and Recovery
  - *Leverage NVM's ability to support fast random writes*
  - *Enable instantaneous recovery from failures*
- Access Methods
  - *Read and write latencies of NVM are asymmetric*
  - *Write-limited access methods such as B+tree*

UC San Diego



UNIVERSITY OF  
TORONTO

# #3: EXECUTION ENGINE


---

- Plan Executor
  - *Write-limited sorting algorithm*
  - *Makes use of selection sort which takes multiple read passes*
- Query Optimizer
  - *Differentiate between reads and writes in cost model*
  - *Factor in byte-addressability of NVM*






**DBMS  
OVERVIEW**



**LOGGING &  
RECOVERY**



**DATA  
PLACEMENT**

# WRITE-AHEAD LOGGING





# WRITE-BEHIND LOGGING

---

- Write-ahead log serves two purposes
  - *Transform random database writes into sequential log writes*
  - *Support transaction rollback*
- NVM supports fast random writes
  - *Directly write data to the multi-versioned database*
  - *Later, record metadata about committed txns in log*




WRITE-BEHIND LOGGING  
VLDB 2016



LET'S TALK ABOUT STORAGE AND RECOVERY METHODS FOR  
NON-VOLATILE MEMORY DATABASE SYSTEMS  
SIGMOD 2015

# WRITE-BEHIND LOGGING



# METADATA FOR INSTANT RECOVERY

- Record failed group commit timestamp gap in log
  - Use it to ignore effects of uncommitted transactions*

*Write-behind logging enables instant recovery  
and avoids data duplication*

$(l_1, l_2)$

$(l_1, l_2)$

List of gaps


Garbage Collection

# EVALUATION

---


- Compare logging protocols in Peloton
  - *Write-Ahead logging*
  - *Write-Behind logging*
- TPC-C benchmark
- Storage devices
  - *Solid-state drive*
  - *Non-volatile memory*

# RECOVERY TIME




# THROUGHPUT

Write-Ahead Logging      Write-Behind Logging






**DBMS  
OVERVIEW**



**LOGGING &  
RECOVERY**



**DATA  
PLACEMENT**

# NVM-AWARE DATA PLACEMENT


---

- Support analytics on a multi-tier storage hierarchy
  - *Cost of first-generation NVM devices*
  - *DRAM + NVM + SSD*



*When should the DBMS migrate data between devices in storage hierarchy?*

# THREE-TIER STORAGE HIERARCHY




# DATA PLACEMENT

---

- Can directly read data from NVM
  - *No need to copy data over to DRAM for reading*
- Cache hot data in DRAM
- Dynamically migrate cold data to SSD
  - *And bring back warm data to NVM*




# THREE-TIER STORAGE HIERARCHY






**DBMS  
OVERVIEW**



**LOGGING &  
RECOVERY**



**DATA  
PLACEMENT**

# THE HOME STRETCH

---

- #1: NVM-aware B+tree (with Microsoft Research)
  - *Write-limited design for NVM*
- #2: Data placement in multi-tier storage hierarchy
  - *Data migration policies*
- #3: Replication
  - *NVM Express over Fabrics library*



**PELOTON**

<http://pelotondb.org>



NVM Ready



Autonomous



Apache Licensed

# END

@joy\_arulraj