

UpBit: Scalable In-Memory Updatable Bitmap Indexing

Manos Athanassoulis

Zheng Yan*

Stratos Idreos



HARVARD
UNIVERSITY

*during an internship at DASlab

Indexing for Analytical Workloads

Column A

30
20
30
10
20
10
30
20

A=10

0
0
0
1
0

A=20

0
1
0
0
1

A=30

1
0
1
0
0

Specialized indexing

- Compact representation of query result
- Query result is readily available

Bitvectors

- Can leverage fast Boolean operators
- Bitwise AND/OR/NOT faster than looping over meta data

Bitmap Indexing Limitations

Column A

30
20
30
10
30
20
10
30
20

A=10

0
0
0
1
0

A=20

0
1
0
0
0

A=30

1
0
1
0
0

Index Size

 Space-inefficient for domains with large cardinality

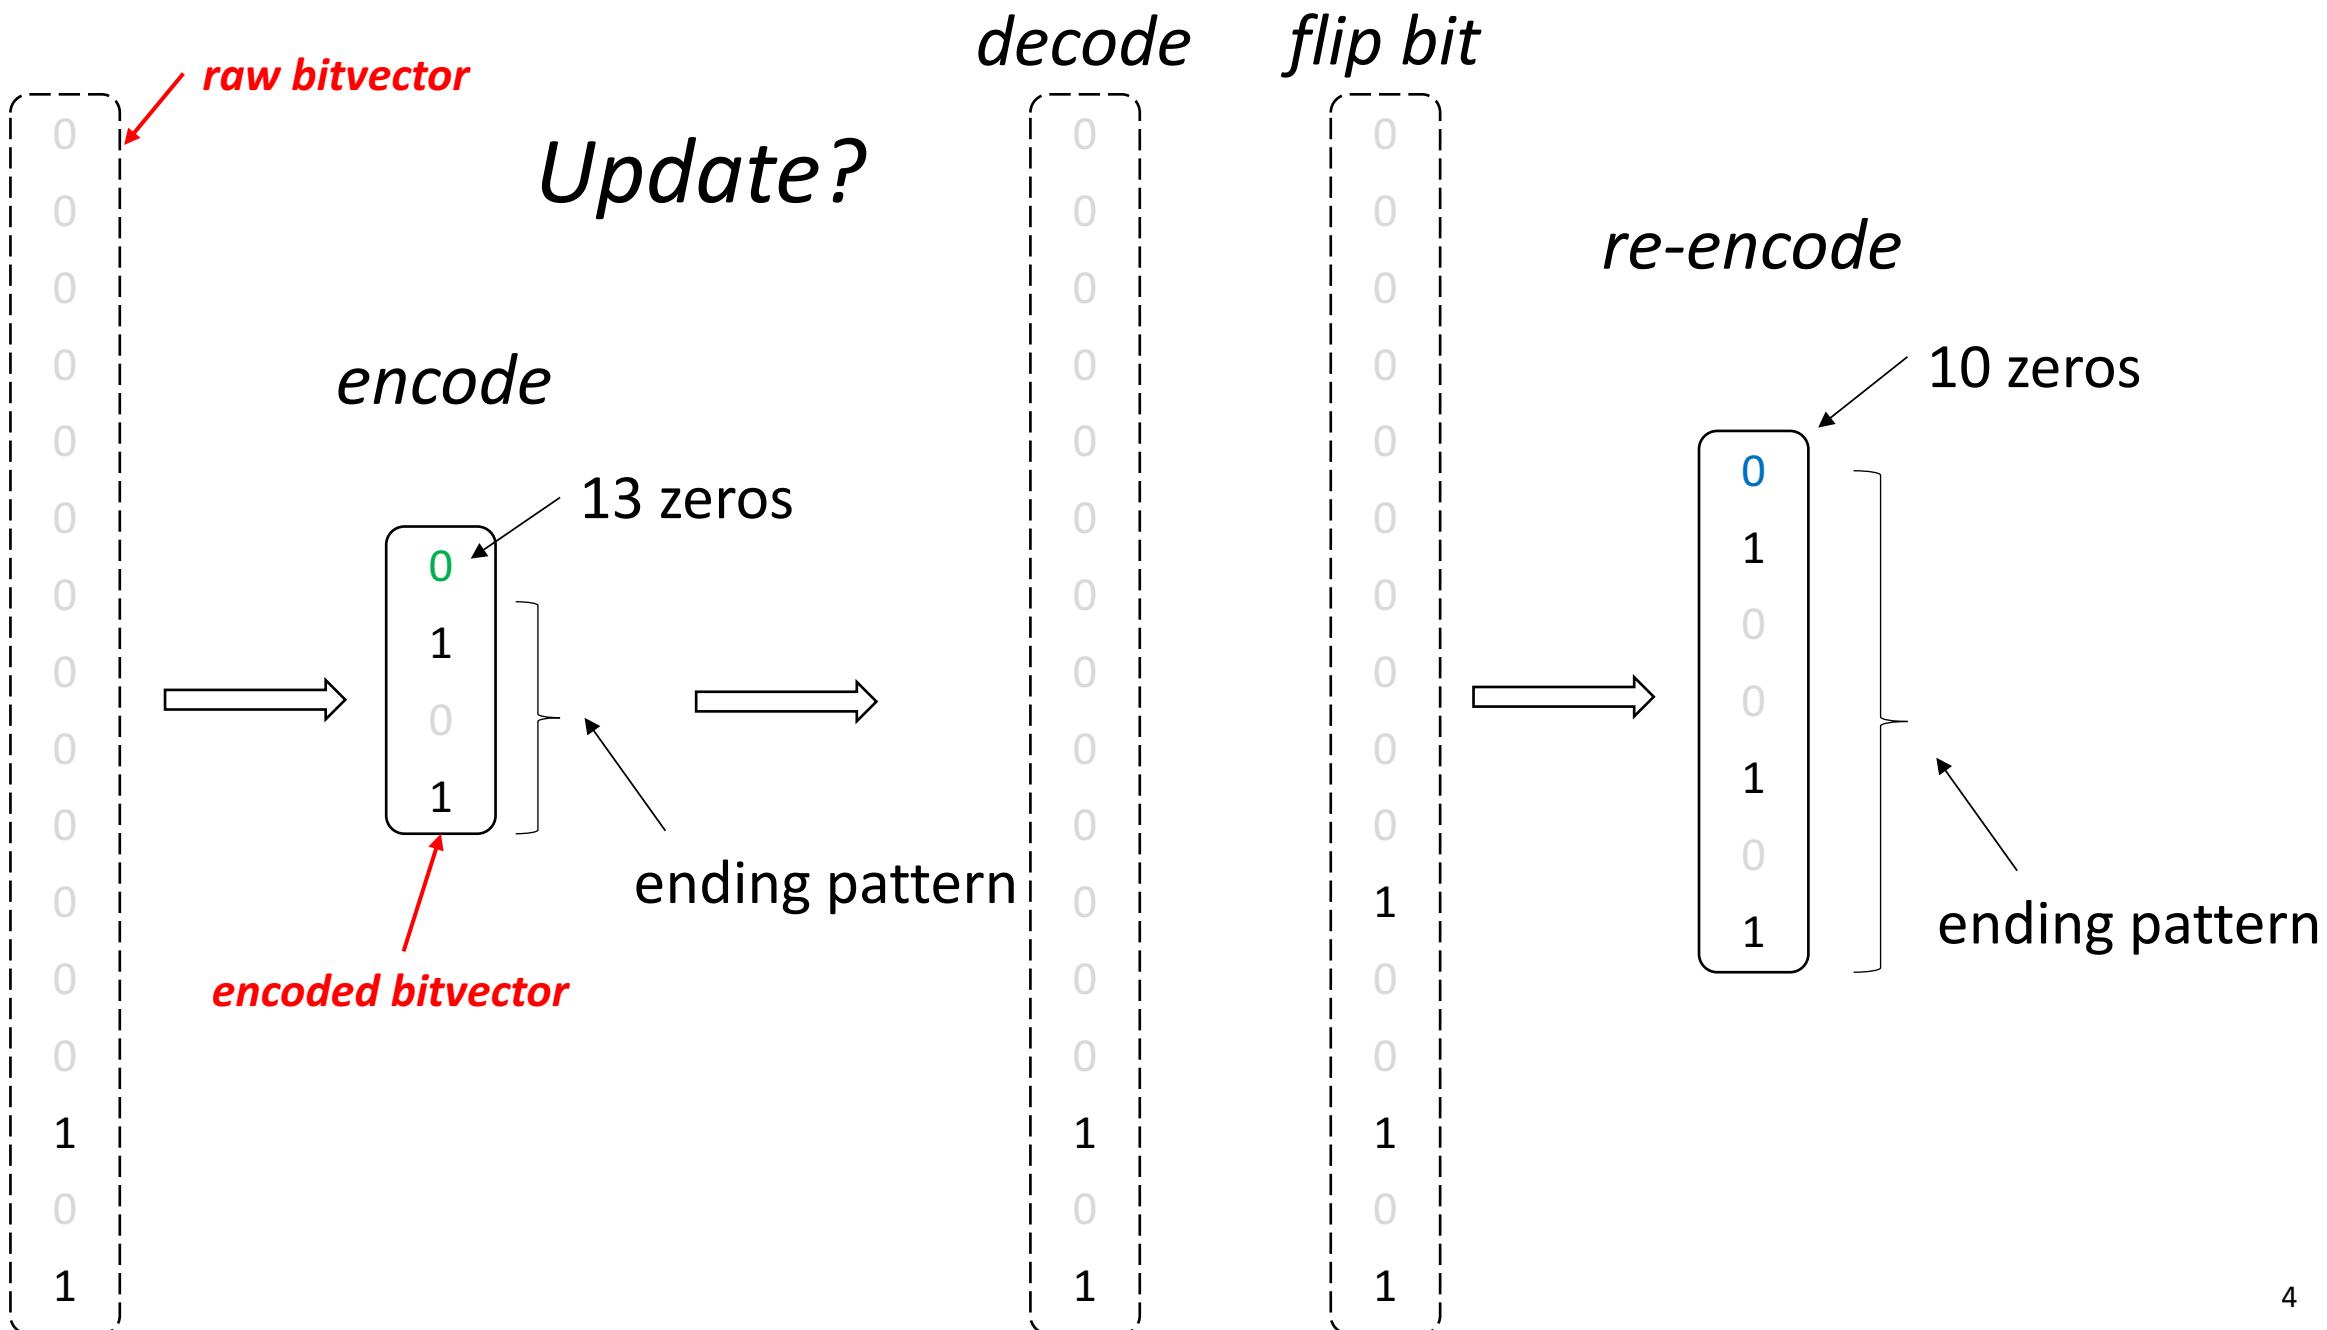
 Addressed by bitvector encoding/compression

core idea: *run-length encoding* in prior work

encoded bitvectors

but ...

 Updating encoded bitvectors is **very** inefficient



Goal

Bitmap Indexing with efficient *Reads & Updates*

Prior Work: Bitmap Indexing and Deletes

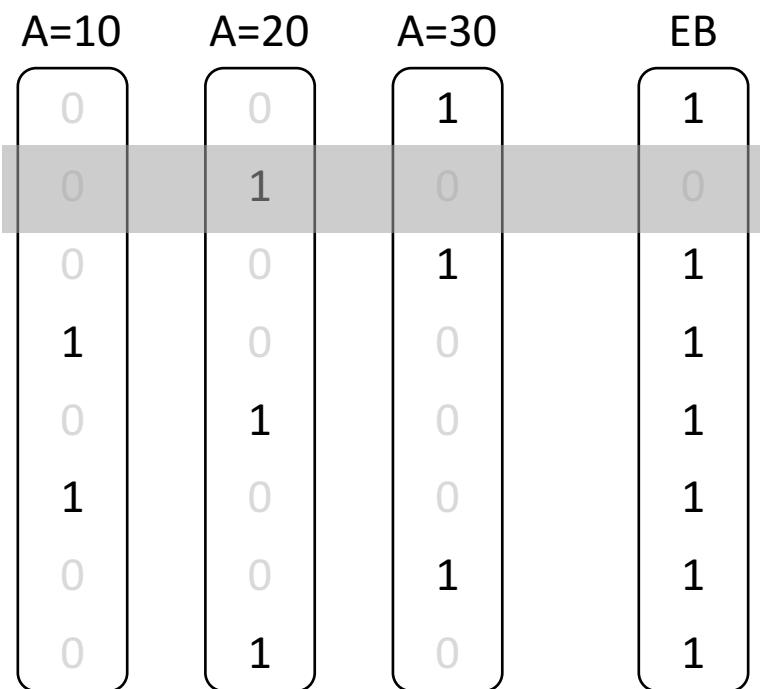
Update Conscious Bitmaps (UCB), SSDBM 2007

A=10	A=20	A=30	EB
0	0	1	1
0	1	0	1
0	0	1	1
1	0	0	1
0	1	0	1
1	0	0	1
0	0	1	1
0	1	0	1

efficient deletes by invalidation
existence bitvector (EB)

Prior Work: Bitmap Indexing and Deletes

Update Conscious Bitmaps (UCB), SSDBM 2007



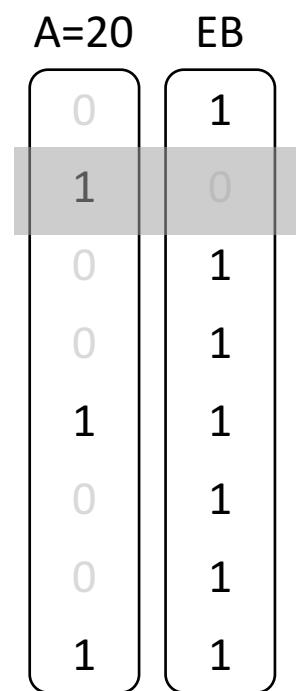
efficient deletes by invalidation
existence bitvector (EB)

reads?

bitwise AND with EB

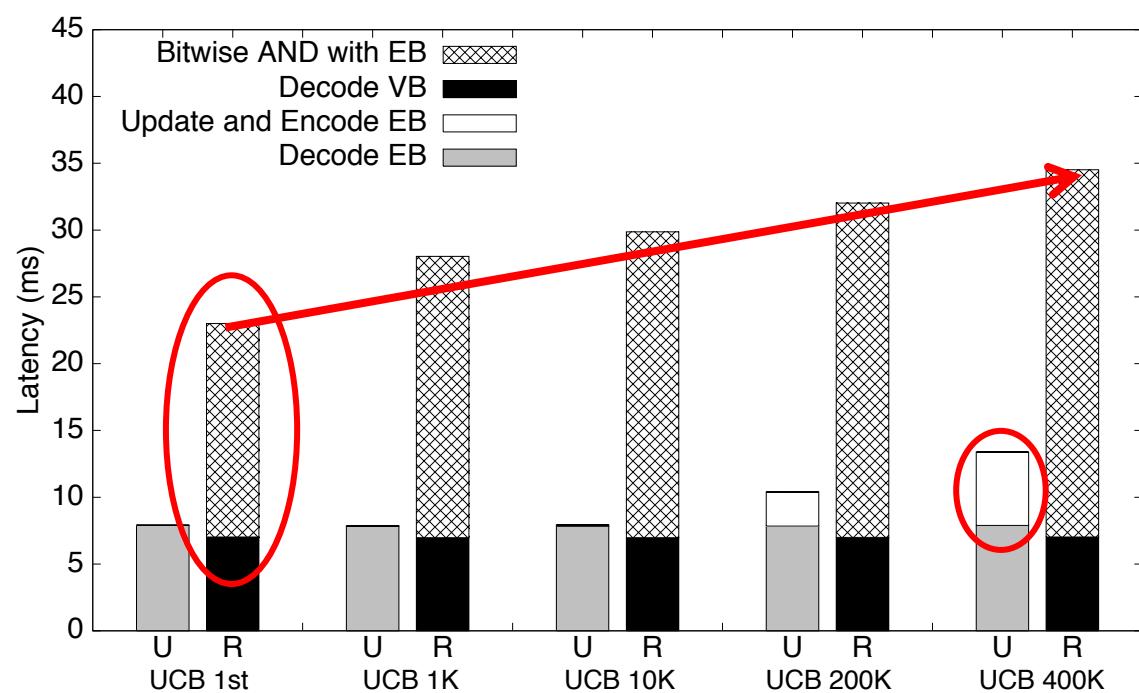
updates?

delete-then-append



Prior Work: Limitations

n=100M tuples, d=100 domain values, 50% updates / 50% reads



read cost increases with #updates

why?

bitwise AND with EB is the bottleneck

update EB is costly for >> #updates

UCB performance does not scale with #updates

single auxiliary bitvector

repetitive bitwise operations

Bitmap Indexing for Reads & Updates



efficient random accesses in compressed bitvectors



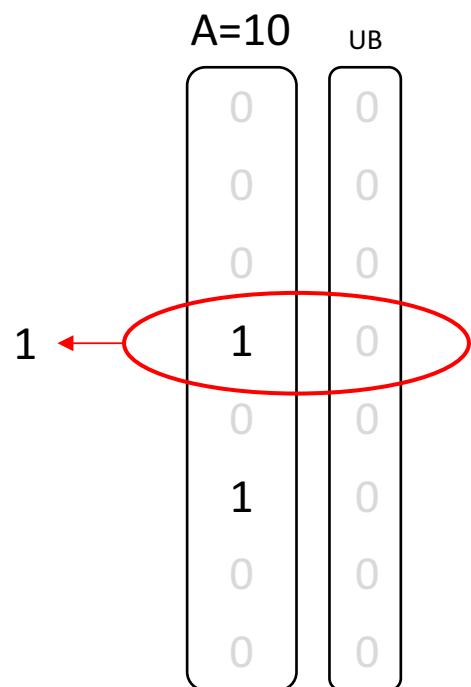
query-driven re-use results of bitwise operations



distribute update cost



Design Element 1: update bitvectors

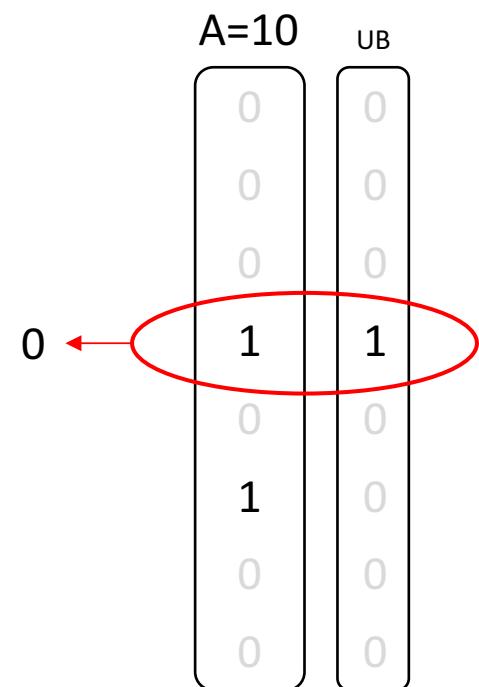


one per value of the domain
initialized to 0s

the current value is the XOR
every update flips a bit on UB



Design Element 1: update bitvectors



one per value of the domain
initialized to 0s

the current value is the XOR
every update flips a bit on UB

... distribute the update burden

Updating UpBit ...

... row 5 to 10

A=10	UB	A=20	UB	A=30	UB
0	0	0	0	1	0
0	0	1	0	0	0
0	0	0	0	1	0
1	0	0	0	0	0
0	0	1	0	0	0
1	0	0	0	0	0
0	0	0	0	1	0
0	0	1	0	0	0

Updating UpBit ...

... row 5 to 10

1. find old value of row 5 (A=20)

A=10	UB	A=20	UB	A=30	UB
0	0	0	0	1	0
0	0	1	0	0	0
0	0	0	0	1	0
1	0	0	0	0	0
0	0	1	0	0	0
1	0	0	0	0	0
0	0	0	0	1	0
0	0	1	0	0	0

Updating UpBit ...

... row 5 to 10

1. find old value of row 5 (A=20)

A=10	UB	A=20	UB	A=30	UB
0	0	0	0	1	0
0	0	1	0	0	0
0	0	0	0	1	0
1	0	0	0	0	0
0	0	1	0	0	0
1	0	0	0	0	0
0	0	0	0	1	0
0	0	1	0	0	0

Updating UpBit ...

A=10	UB	A=20	UB	A=30	UB
0	0	0	0	1	0
0	0	1	0	0	0
0	0	0	0	1	0
1	0	1	1	0	0
0	0	0	0	0	0
1	0	0	0	0	0
1	0	0	0	0	0
0	0	0	0	1	0
0	0	1	0	0	0

... row 5 to 10

1. find old value of row 5 (A=20)
2. flip bit of row 5 of UB of A=20

Updating UpBit ...

A=10	UB	A=20	UB	A=30	UB
0	0	0	0	1	0
0	0	1	0	0	0
0	0	0	0	1	0
1	0	0	0	0	0
0	1	1	1	0	0
1	0	0	0	0	0
0	0	0	0	0	0
0	0	1	0	1	0

... row 5 to 10

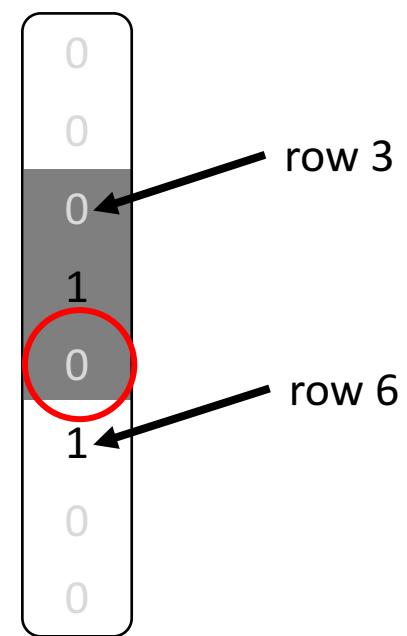
1. find old value of row 5 ($A=20$)
2. flip bit of row 5 of UB of $A=20$
3. flip bit of row 5 of UB of $A=10$

can we speed up step 1?



Design Element 2: fence pointers

efficient access of compressed bitvectors
fence pointers



we can find row 5 without decoding & scanning the whole bitvector

Updating UpBit ...

... row 5 to 10

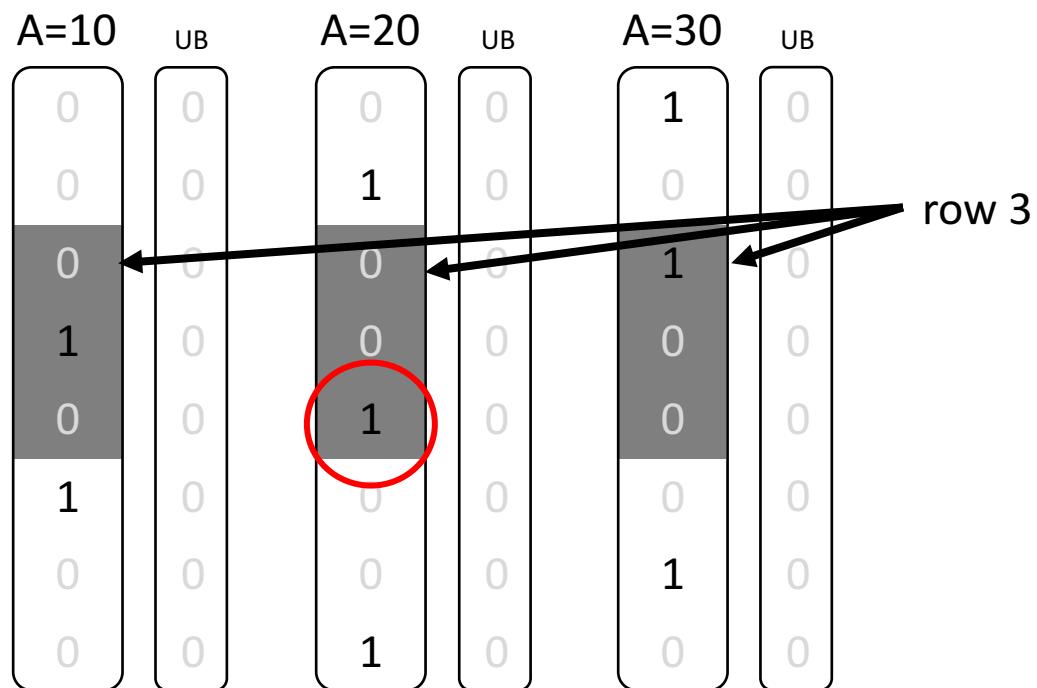
1. find old value of row 5 (A=20)

A=10	UB	A=20	UB	A=30	UB
0	0	0	0	1	0
0	0	1	0	0	0
0	0	0	0	1	0
1	0	0	0	0	0
0	0	1	0	0	0
1	0	0	0	0	0
0	0	0	0	1	0
0	0	1	0	0	0

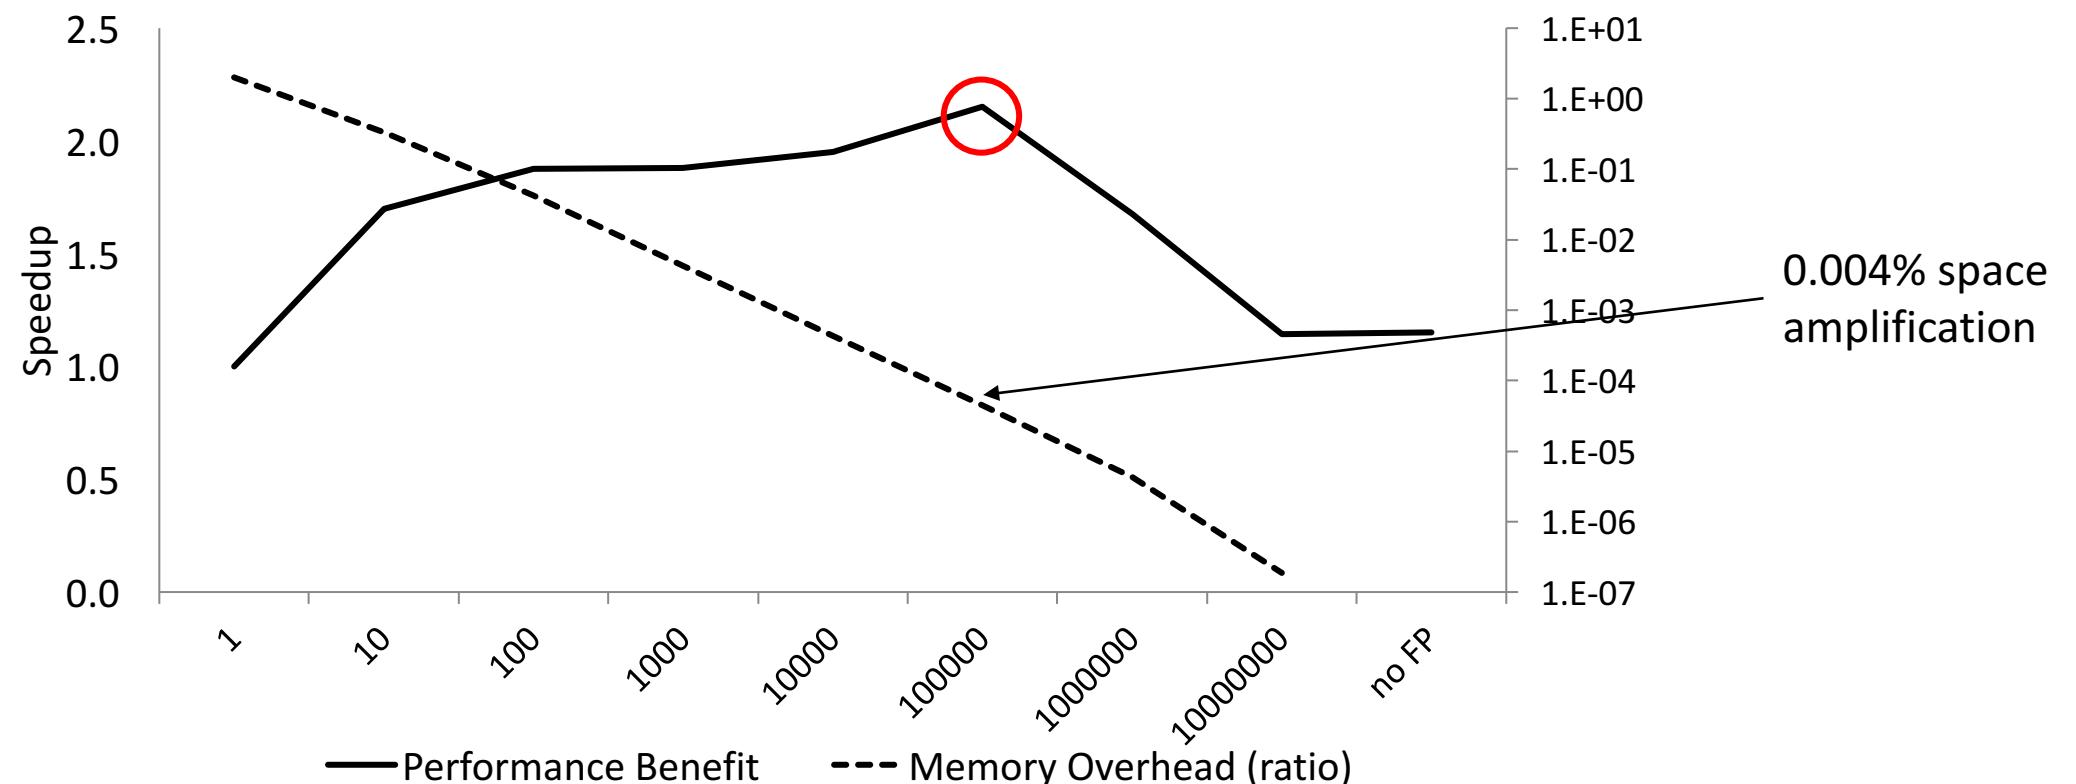
Updating UpBit (with fence pointers)...

... row 2 to 10

1. find old value of row 2 ($A=20$)
using fence pointers



How dense should the *fence pointers* be?



fence pointers every 10^5 entries

Querying

Querying UpBit ...

... A = 20

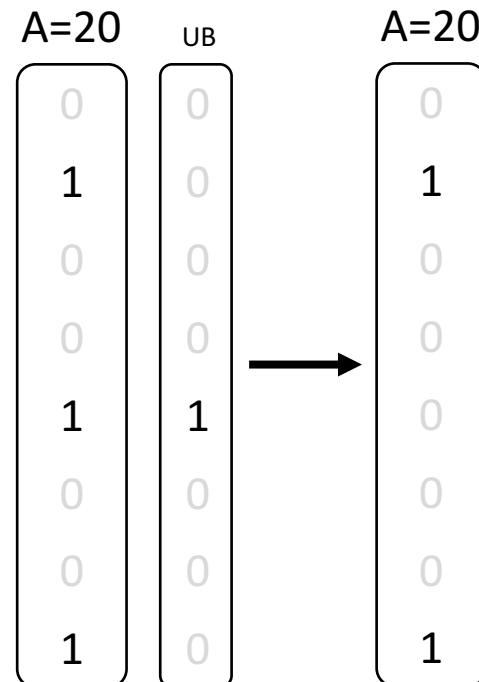
Return the XOR of A=20 and UB

A=10	UB	A=20	UB	A=30	UB
0	0	0	0	1	0
0	0	1	0	0	0
0	0	0	0	1	0
1	0	0	0	0	0
0	1	1	1	0	0
1	0	0	0	0	0
0	0	0	0	0	0
0	0	1	0	1	0

Querying UpBit ...

... A = 20

Return the XOR of A=20 and UB

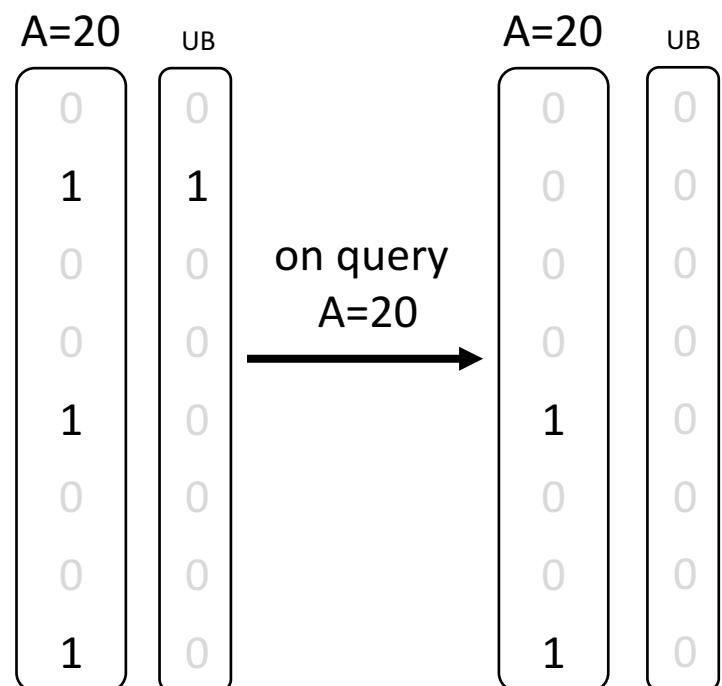


can we re-use the result?

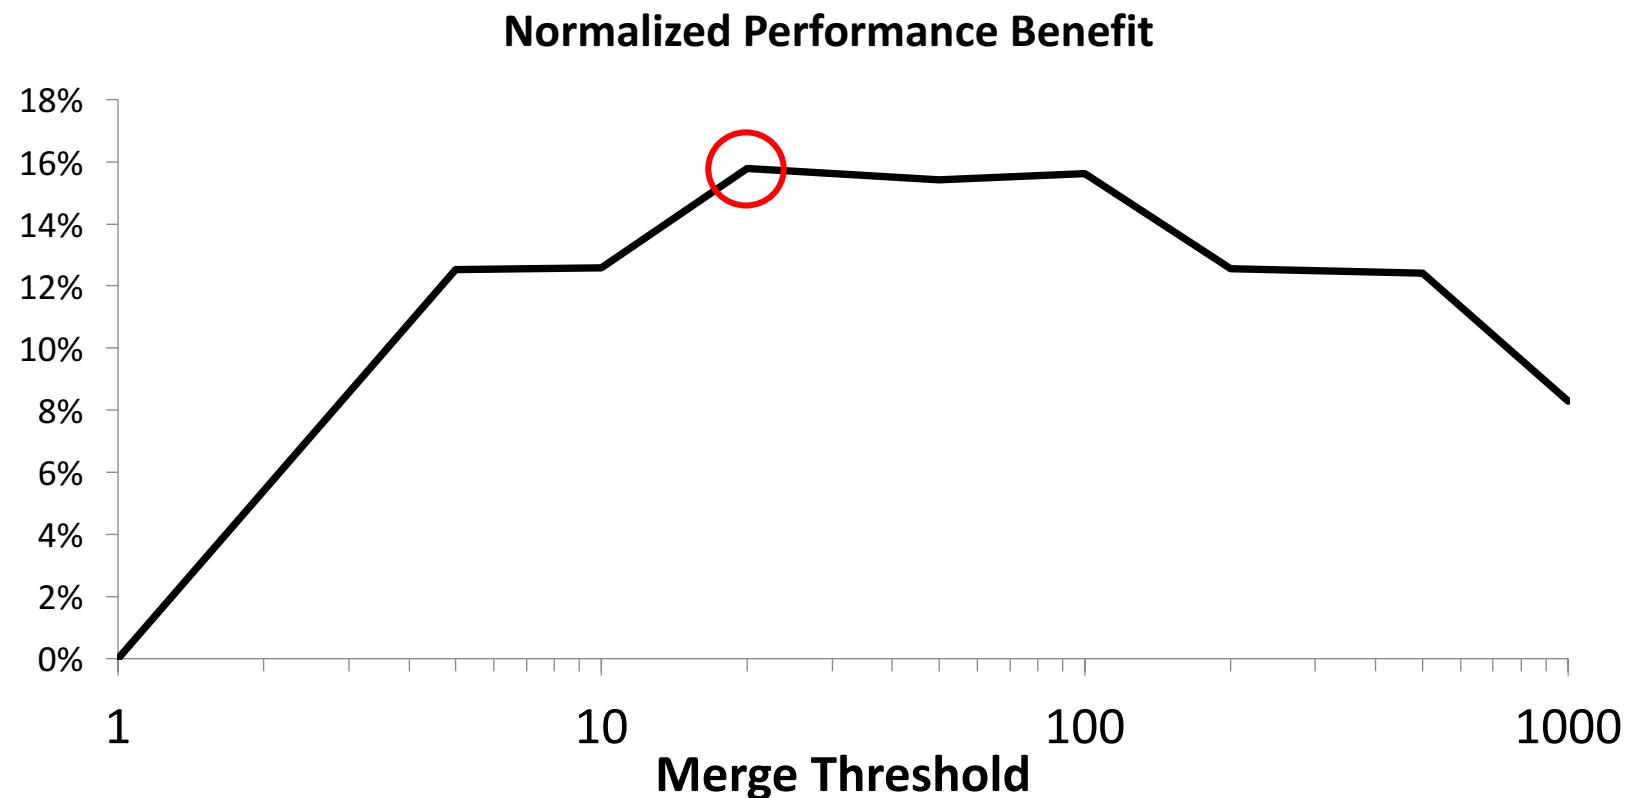


Design Element 3: query-driven merging

maintain high compressibility of UB
query-driven merging



How frequently to merge UB back to VB?



merge back when 20 updates have been recorded

UpBit



update bitvectors



fence pointers



query-driven UB merging

Goals

Faster updates

Faster reads
(than update-optimized)

Fast overall latency

Experiments

Synthetic data

n : # tuples

d : # domain values (cardinality)

q : # queries

u : % updates in the workload

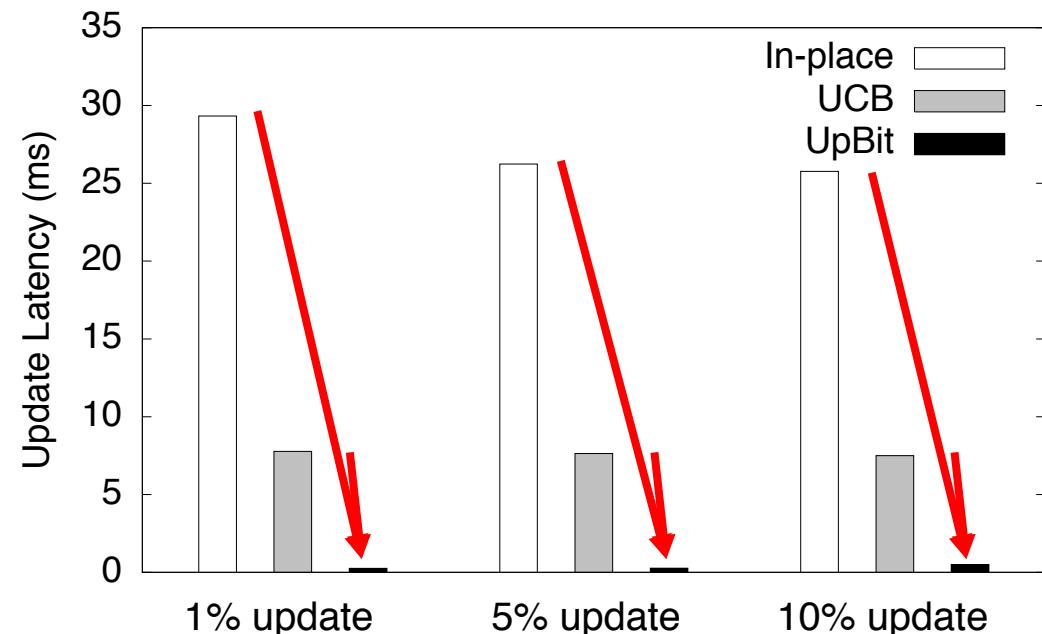
TPC-H SF 100 data

Prototype C++ Implementation of UpBit using FastBit

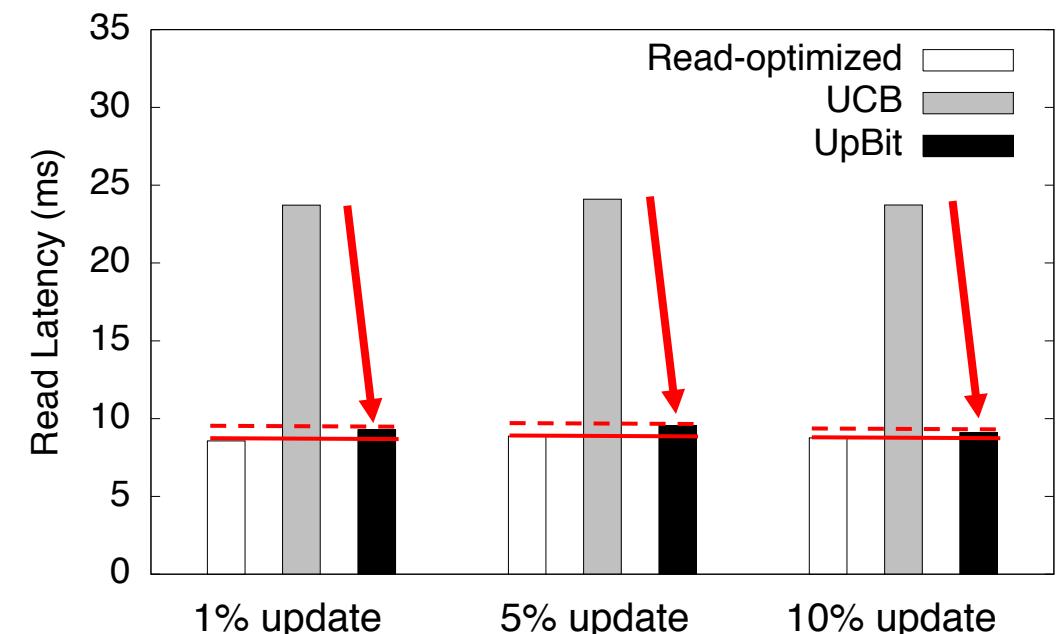
Integrated in a prototype column-store system

UpBit supports very efficient updates

$n=100M$ tuples, $d=100$ domain values
100k queries (varying % of updates)



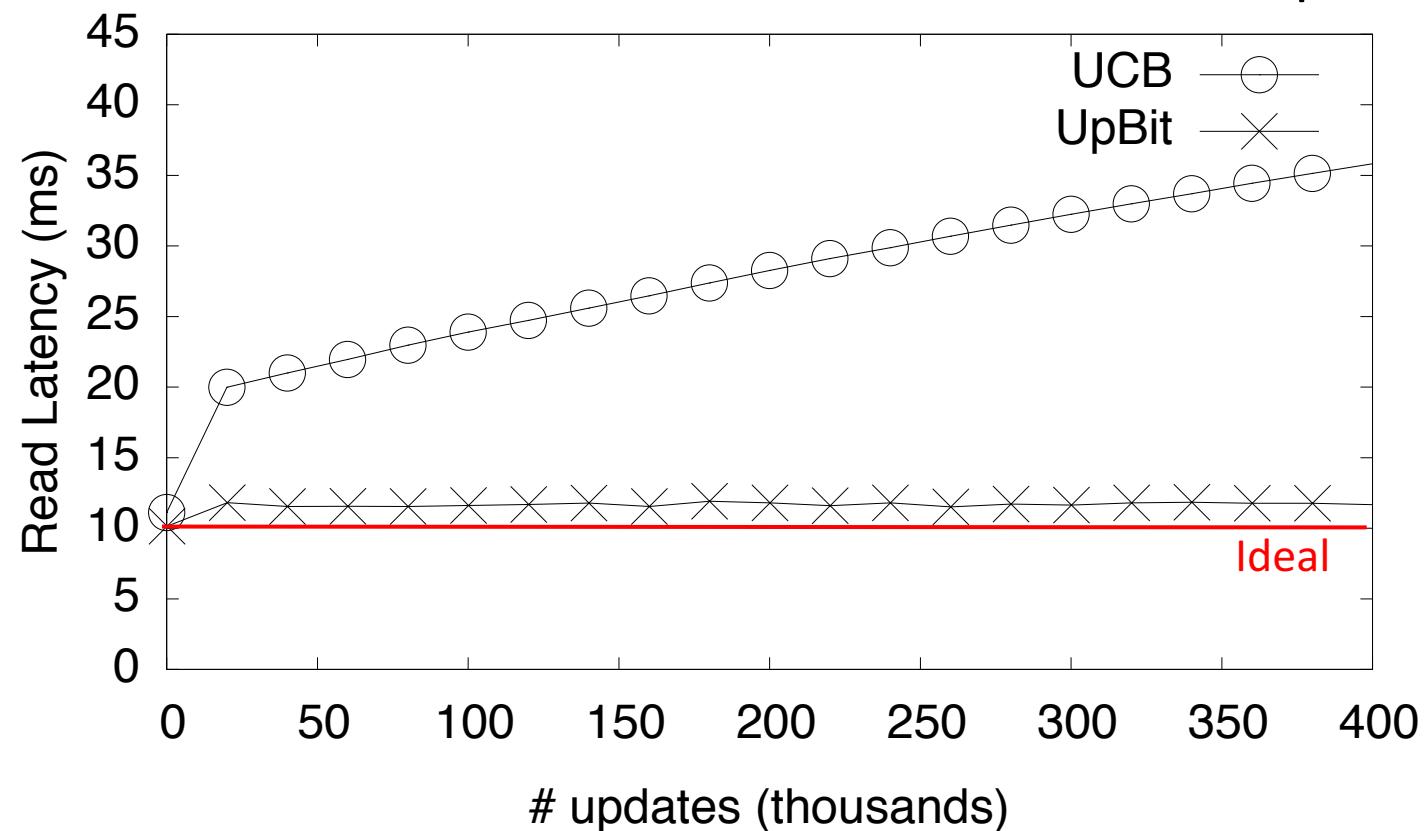
updates: 15-29x faster than UCB
51-115x faster than in-place



only 8% read overhead over optimal
3x faster reads than UCB

UpBit offers robust reads

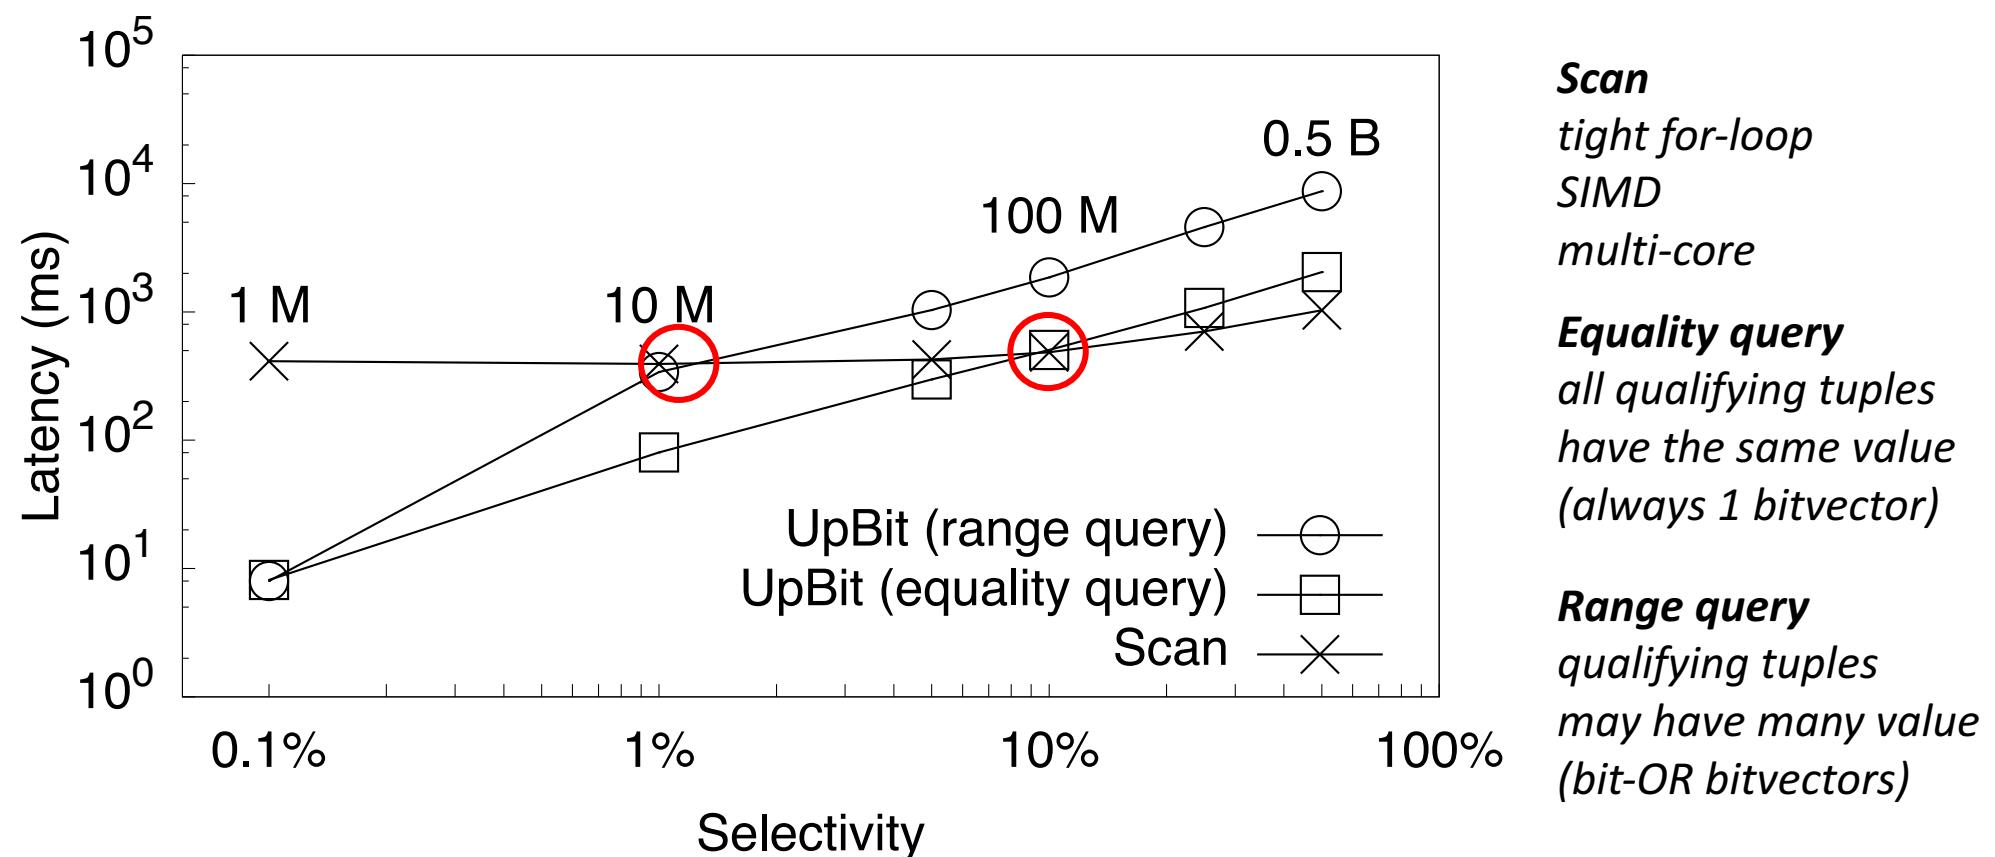
n=100M tuples, d=100 domain values
50%/50% update/read queries



scalable read latency

UpBit as a general index: UpBit vs. Scan

$n = 1B$, $d = 1000$ distinct domain values (range)
 $n = 1B$, d varies for equality: 1000, 100, 10, 1



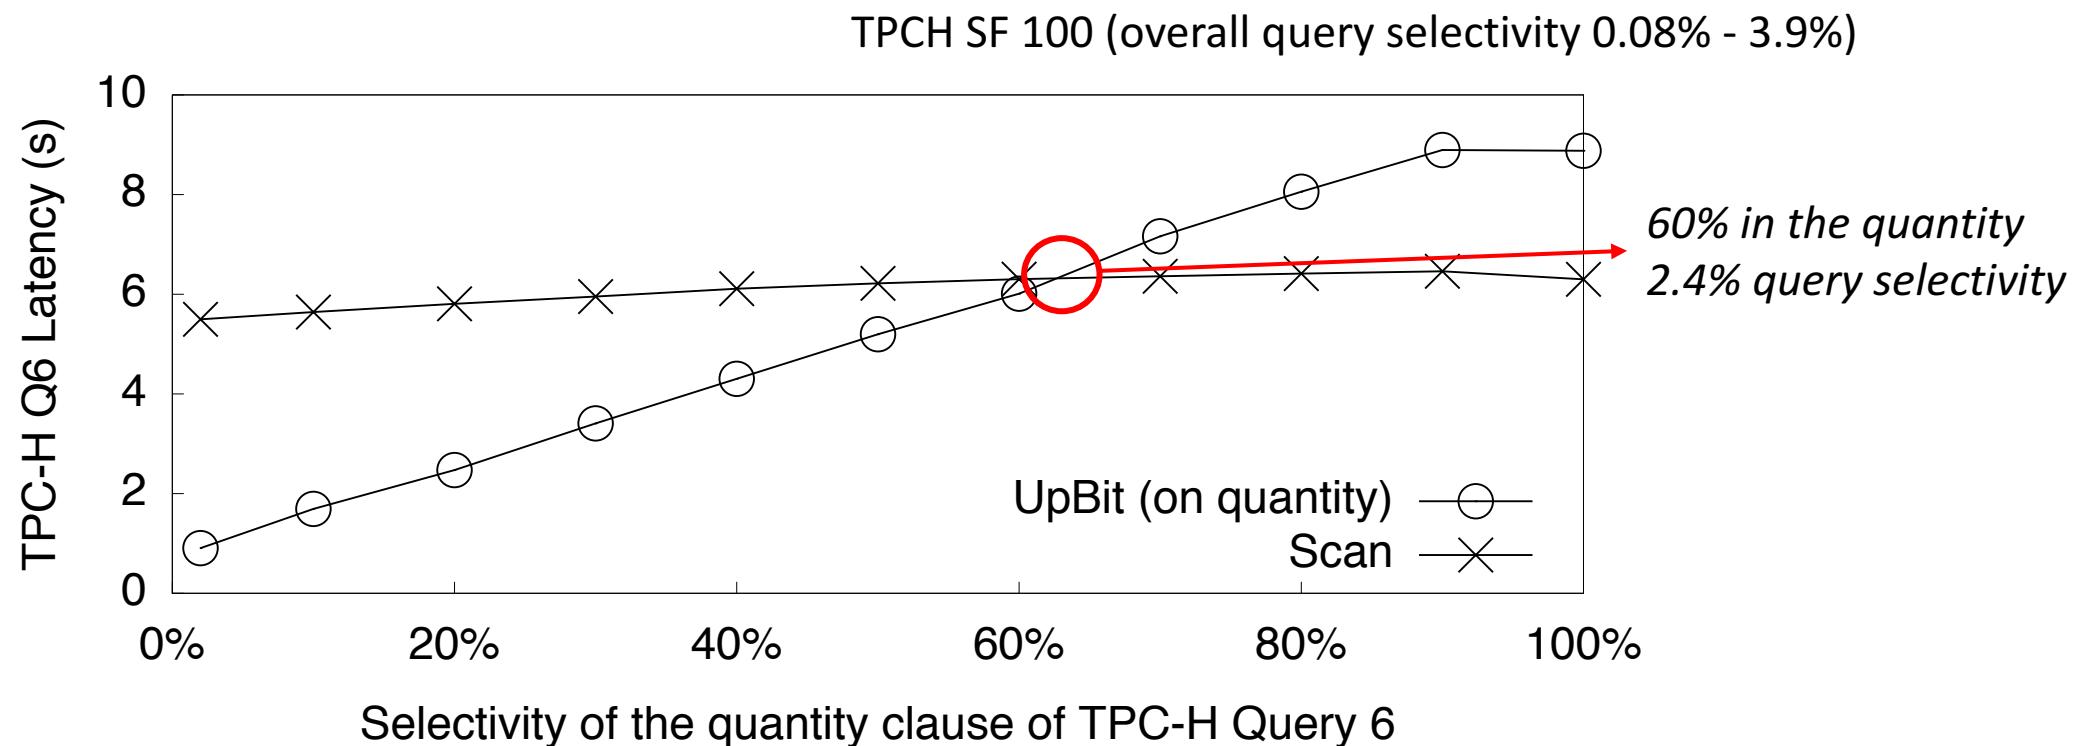
Scan
tight for-loop
SIMD
multi-core

Equality query
all qualifying tuples
have the same value
(always 1 bitvector)

Range query
qualifying tuples
may have many value
(bit-OR bitvectors)

break-even: 10% for equality queries and 1% for range queries

UpBit as a general index: UpBit vs. Scan TPCH Q6



```
SELECT sum(l_extendedprice * l_discount) as revenue
FROM lineitem
WHERE l_shipdate >= date '[DATE]'
AND l_shipdate < date '[DATE]' + interval '1' year
```

an update-aware bitmap index is a viable general-purpose index

UpBit: achieving scalable updates



distribute the update burden
update bitvectors



efficient bitvector accesses
fence pointers



avoid redundant bitwise operations
query-driven merging of UB

Thanks!



DASlab



HARVARD
UNIVERSITY

<http://daslab.seas.harvard.edu/rum/>