



BROWN

REVISING REUSE IN (MAIN-MEMORY) DATABASES

KAYHAN DURSUN*, CARSTEN BINNIG, UGUR CETINTEMEL, TIM KRASKA

REUSE IN THE REAL WORLD ...

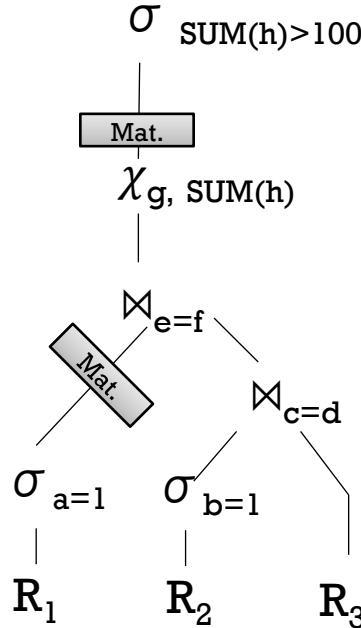


... MIGHT HELP!

... BUT MIGHT ALSO HURT!



REUSE IN DATABASES MIGHT HURT ...

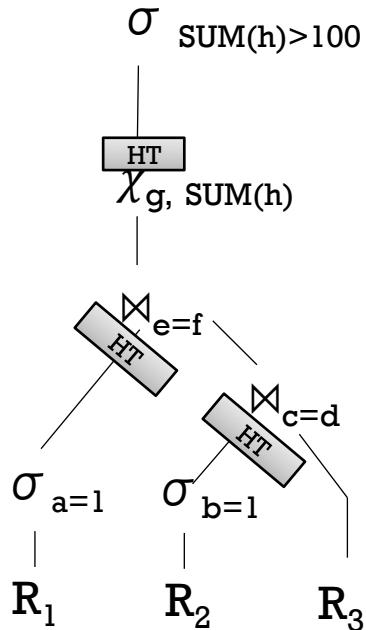


1. **Reason about reuse** potential of intermediate results
2. **Add additional materialization** operations to query plan
3. Analyze if subsequent query can **reuse cached intermediates**



Unclear if extra costs pay off in future

REUSE WITHOUT REGRET ...



Main idea: Reuse internal structures

- Keep internal data structures
- Reuse for subsequent queries

Savings are two-fold

- **No** additional **materialization costs**
- **No** need to **re-create internal structures**

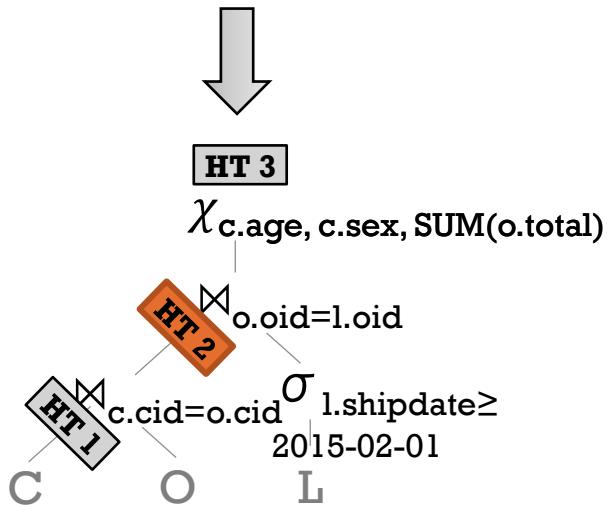
More robust towards different reuse-potentials

This talk: reuse hash tables for joins + aggregations

CASE I: EXACT REUSE

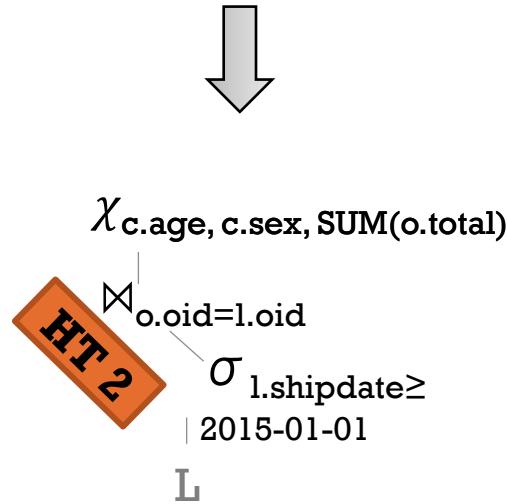
Initial Query

```
SELECT c.age, c.sex,  
       SUM(o.total)  
FROM C, O, L  
WHERE <join-conditions>  
AND l.shipdate≥2015-02-01  
GROUP BY c.age, c.sex
```



Reuse Query

```
SELECT c.age, c.sex,  
       SUM(o.total)  
FROM C, O, L  
WHERE <join-conditions>  
AND l.shipdate≥2015-01-01  
GROUP BY c.age, c.sex
```

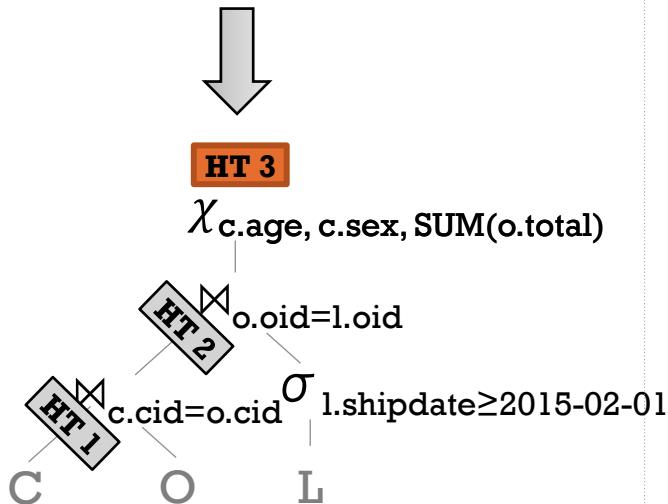


Exact-Reuse HT2

CASE II: PARTIAL REUSE

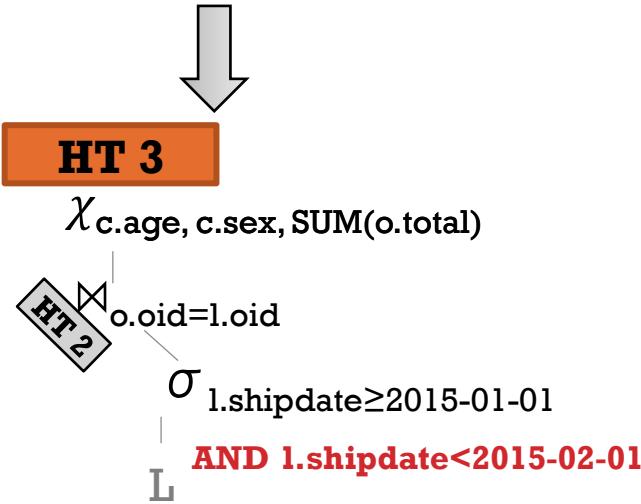
Initial Query

```
SELECT c.age, c.sex,  
       SUM(o.total)  
FROM C, O, L  
WHERE <join-conditions>  
AND l.shipdate≥2015-02-01  
GROUP BY c.age, c.sex
```



Reuse Query

```
SELECT c.age, c.sex,  
       SUM(o.total)  
FROM C, O, L  
WHERE <join-conditions>  
AND l.shipdate≥2015-01-01  
GROUP BY c.age, c.sex
```

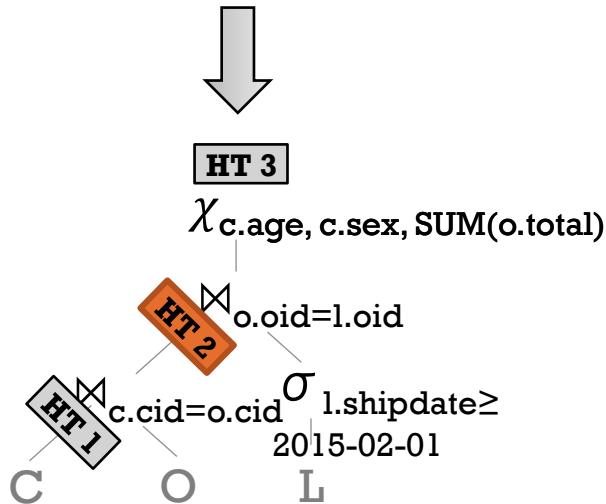


Partial-Reuse HT3:
Add missing tuples

CASE III: SUBSUMING REUSE

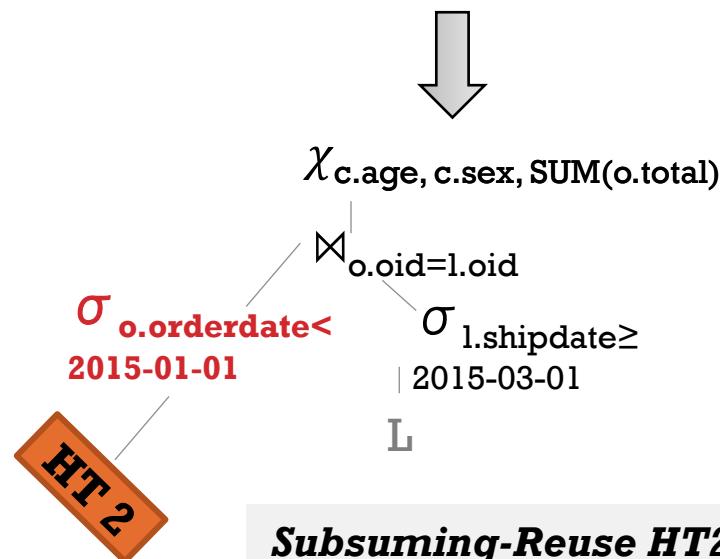
Initial Query

```
SELECT c.age, c.sex,
       SUM(o.total)
  FROM C, O, L
 WHERE <join-conditions>
 AND l.shipdate≥2015-02-01
 GROUP BY c.age, c.sex
```



Reuse Query

```
SELECT c.age, c.sex,
       SUM(o.total)
  FROM C, O, L
 WHERE <join-conditions>
 AND l.shipdate≥2015-02-01
AND o.orderdate<2015-01-01
 GROUP BY c.age, c.sex
```



Subsuming-Reuse HT2:
Filter false positives

HASH-STASH ARCHITECTURE

```
SELECT o.total  
FROM L JOIN O WHERE  
o.orderdate≥2015-01-01
```

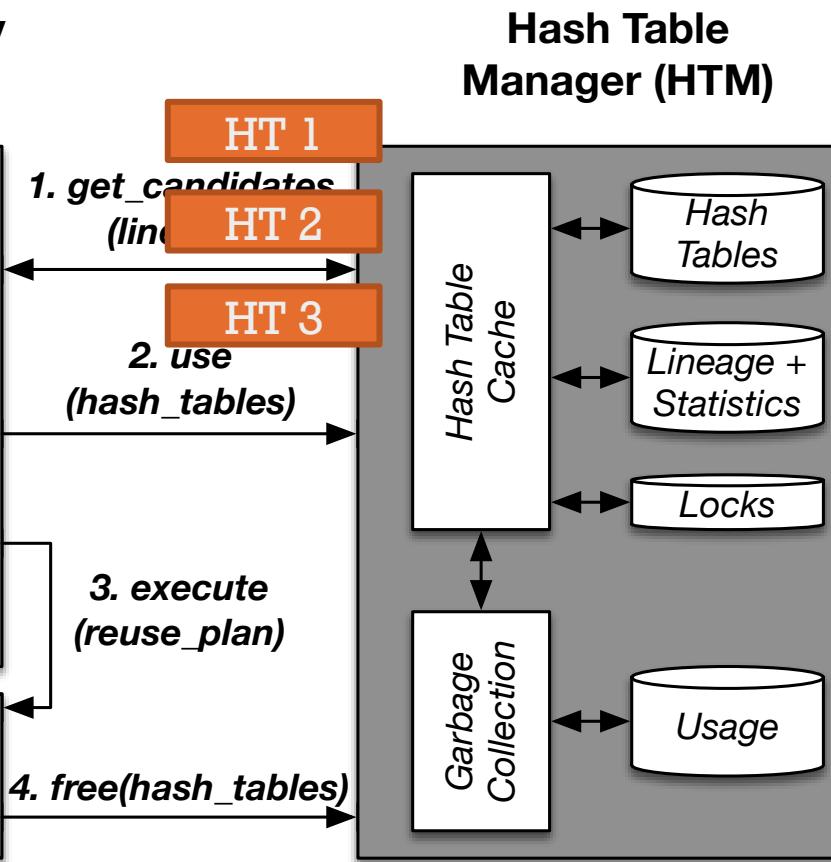
Optimizer (RQO)

Which HT to reuse
for the join?

Rewrite query
and execute it

Executor

DBMS Runtime



COST-MODEL FOR REUSE IN JOINS

Cost of a Hash- \bowtie : $c_{build}(HT) + c_{probe}(HT)$

How is that different for a RHJ?

$$c_{RHJ} = \underline{c_{resize}^R(HT)} + \underline{\underline{c_{build}^R(HT)}} + \underline{\underline{\underline{c_{probe}^R(HT)}}}$$

1. add only
missing tuples

2. higher per-tuple cost
(due to add. tuples)

3. post-filter false
positives (i.e. a=0)

Reuse Query:

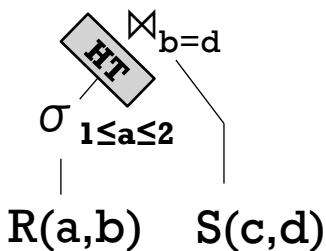


Table R

a	b
0	1
0	2
1	3
1	4
2	5
3	6

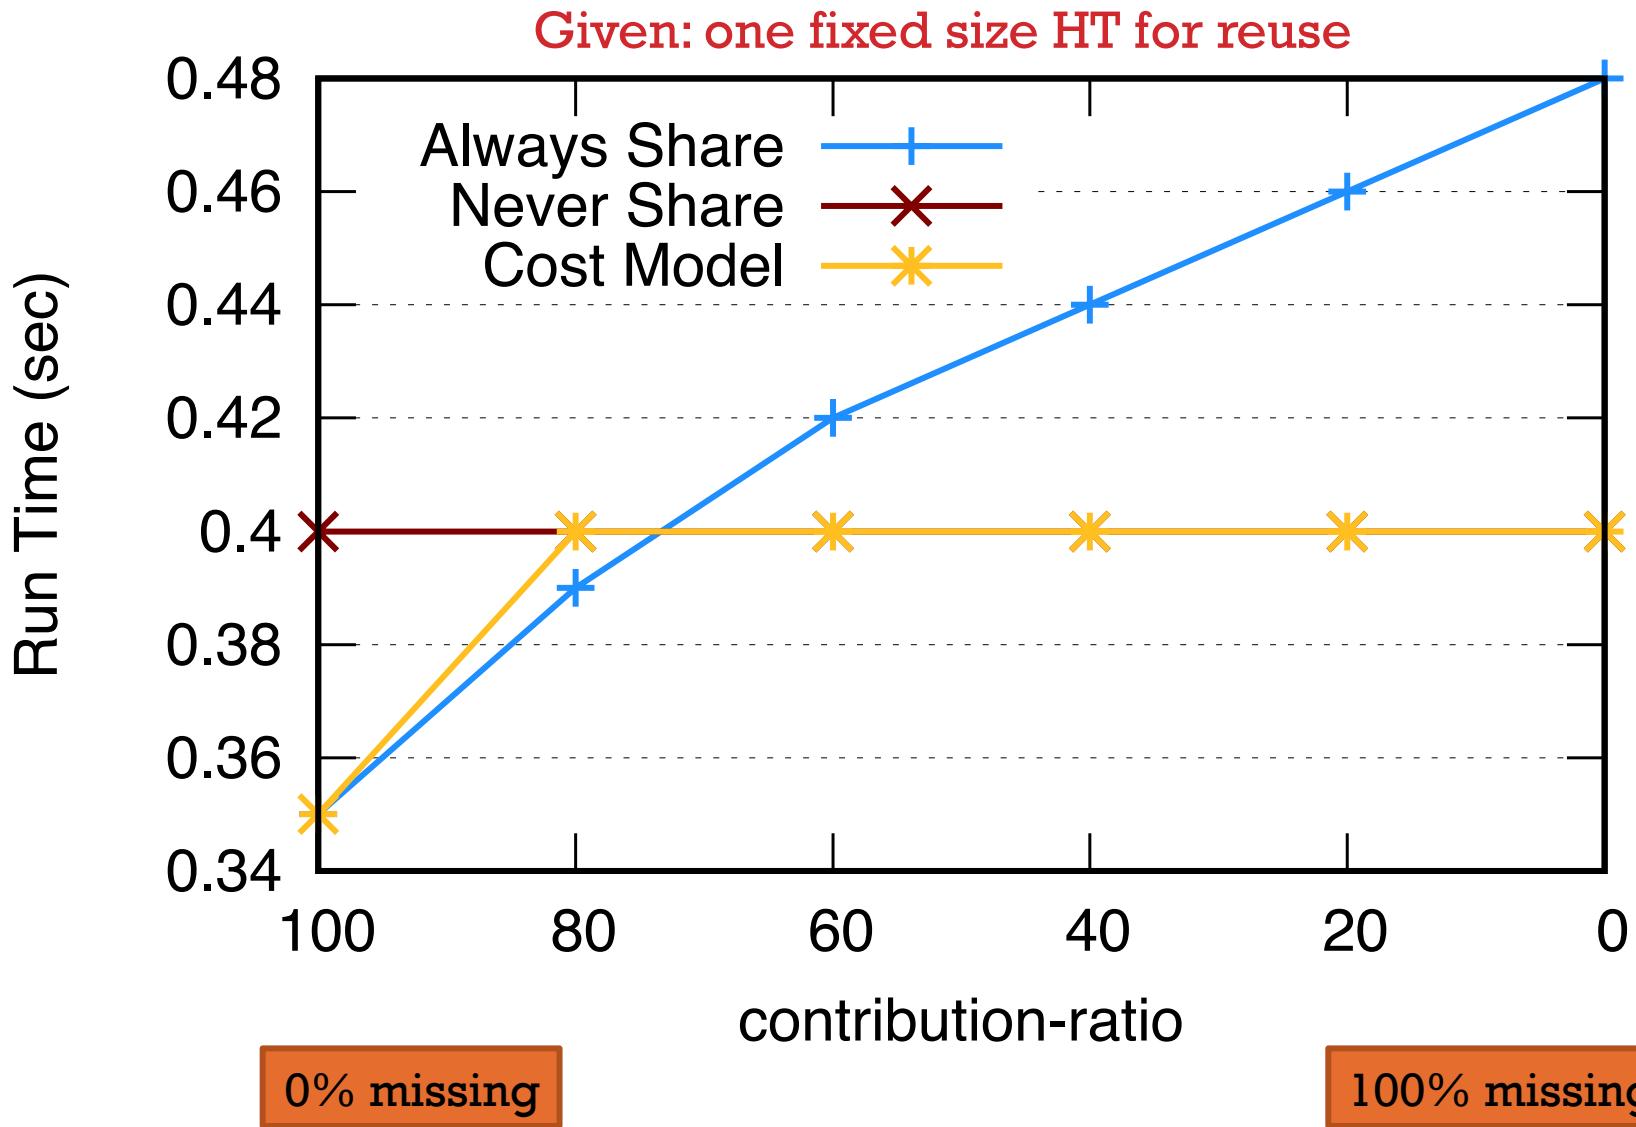
HT: $0 \leq a \leq 1$

b	[a,b]
1	[0,1]
2	[0,2]
3	[1,3]
4	[1,4]
5	[2,5]

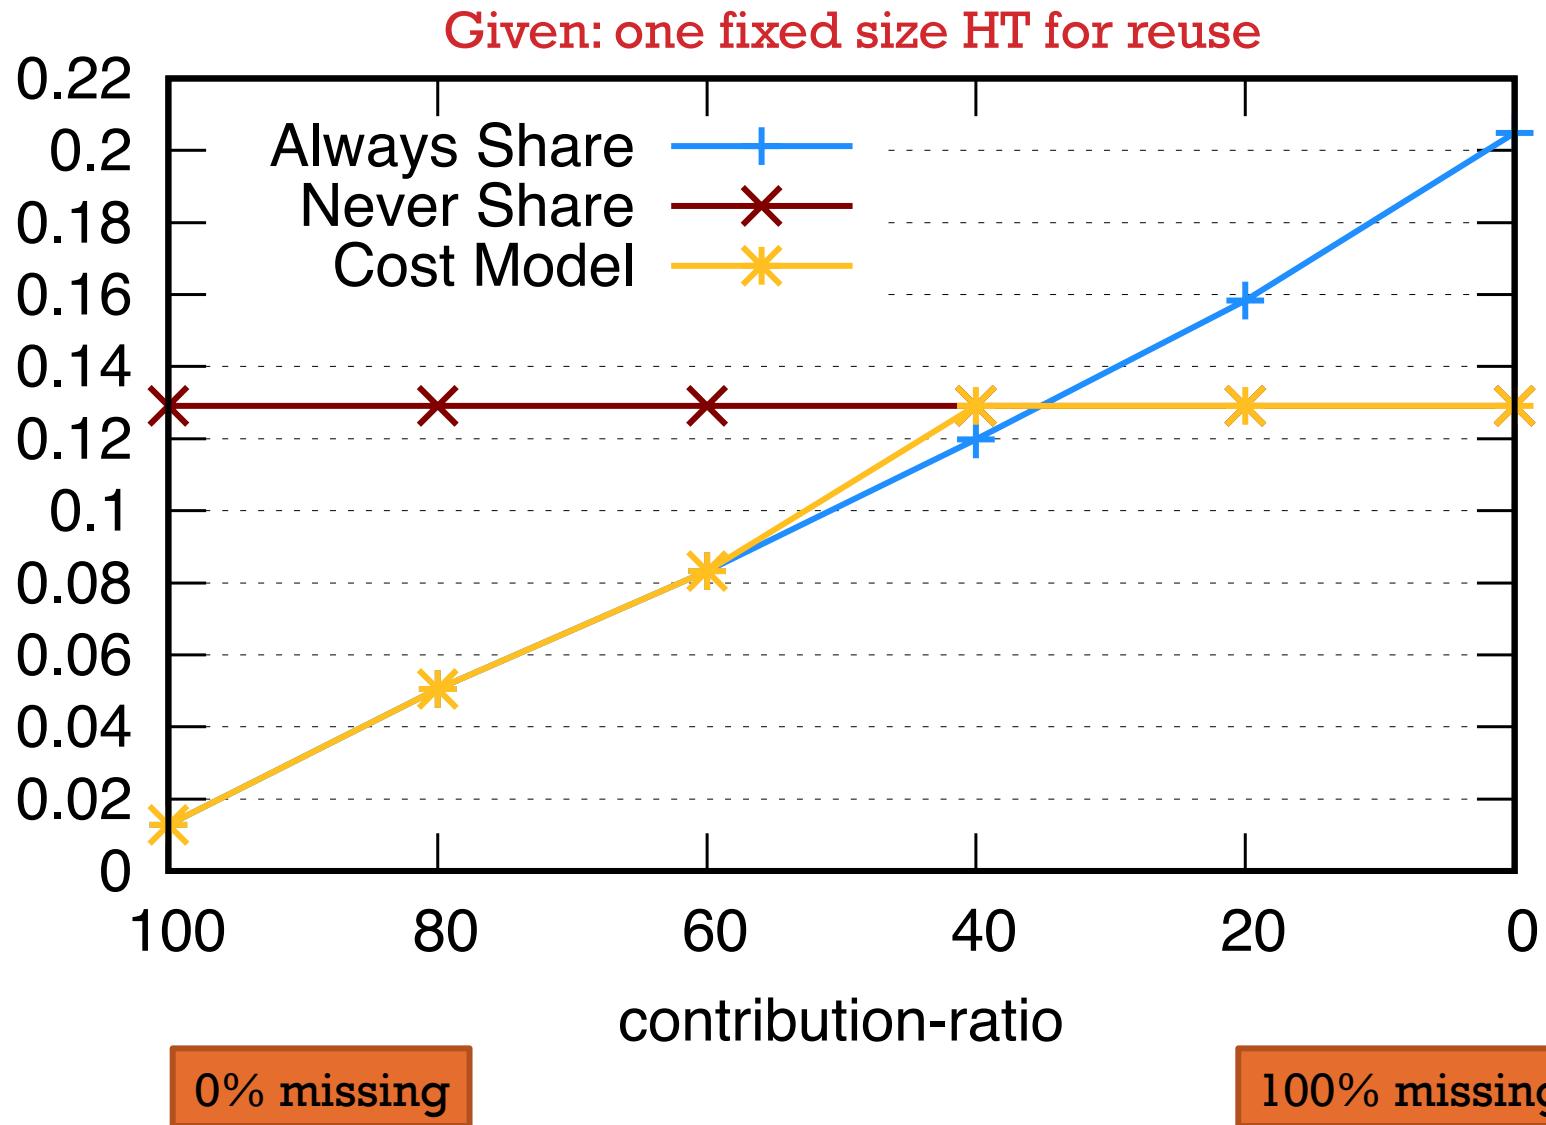
Table S

c	d
1	1
2	1
3	2
4	3
5	4
6	5

JOIN: REUSE OR NOT REUSE?



AGGREGATION: REUSE OR NOT REUSE?



ADDITIONAL OPTIMIZATIONS

Store additional attributes in HTs

- Enables post-filtering false positives
- Only if extension does not exceed cache line

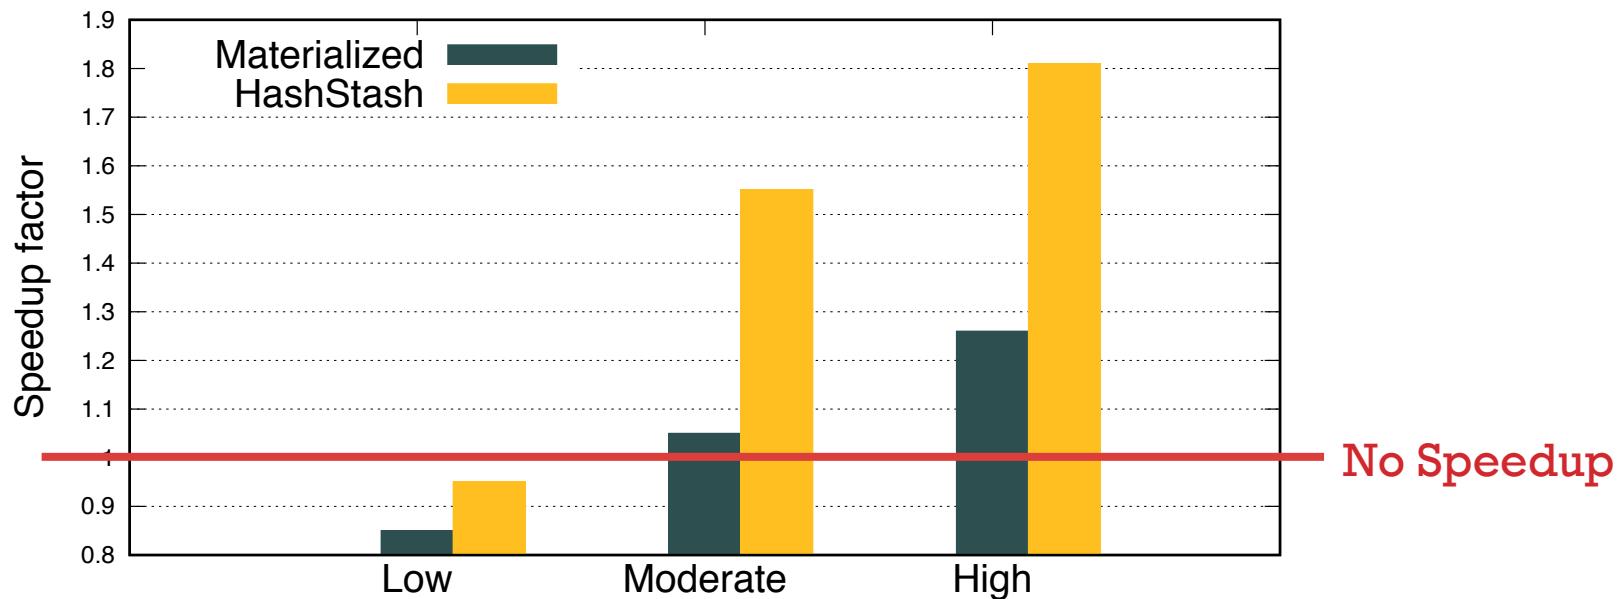
Aggregate Optimizations

- Always store COUNT and SUM instead of AVG
- Add additional aggregate functions

Speculative Optimizations

- Join re-ordering to produce “better” hash-tables
- Store aggregate hash table on finer granularity

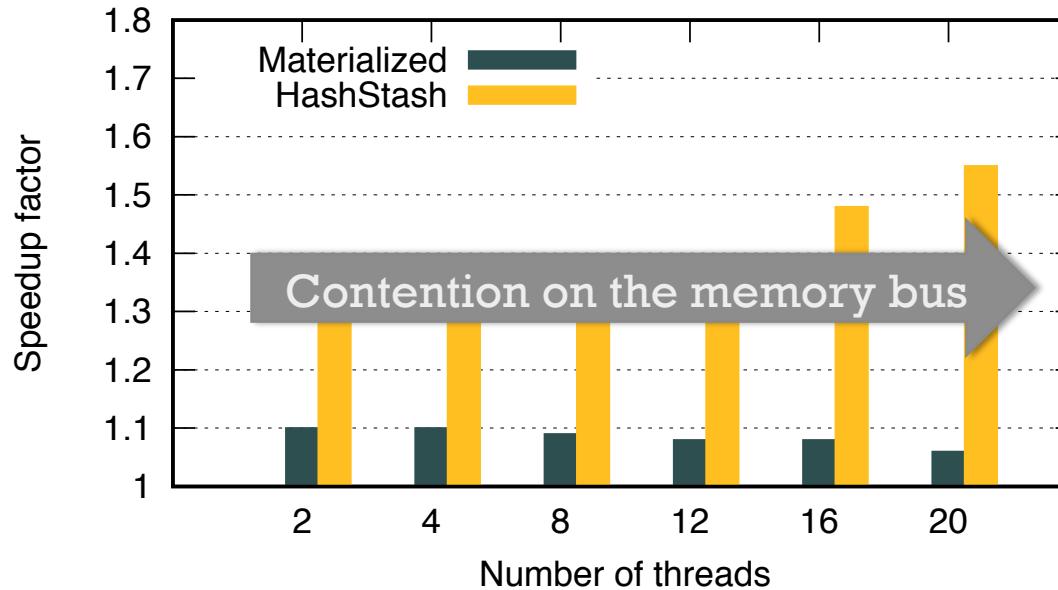
RESULTS: REUSE EFFICIENCY



Workloads: 64 queries with different reuse potentials

- **Low:** 1% of the cached data is reused on average
- **Medium:** 10% of the cached data is reused on average
- **High:** 50% of the cached data is reused on average

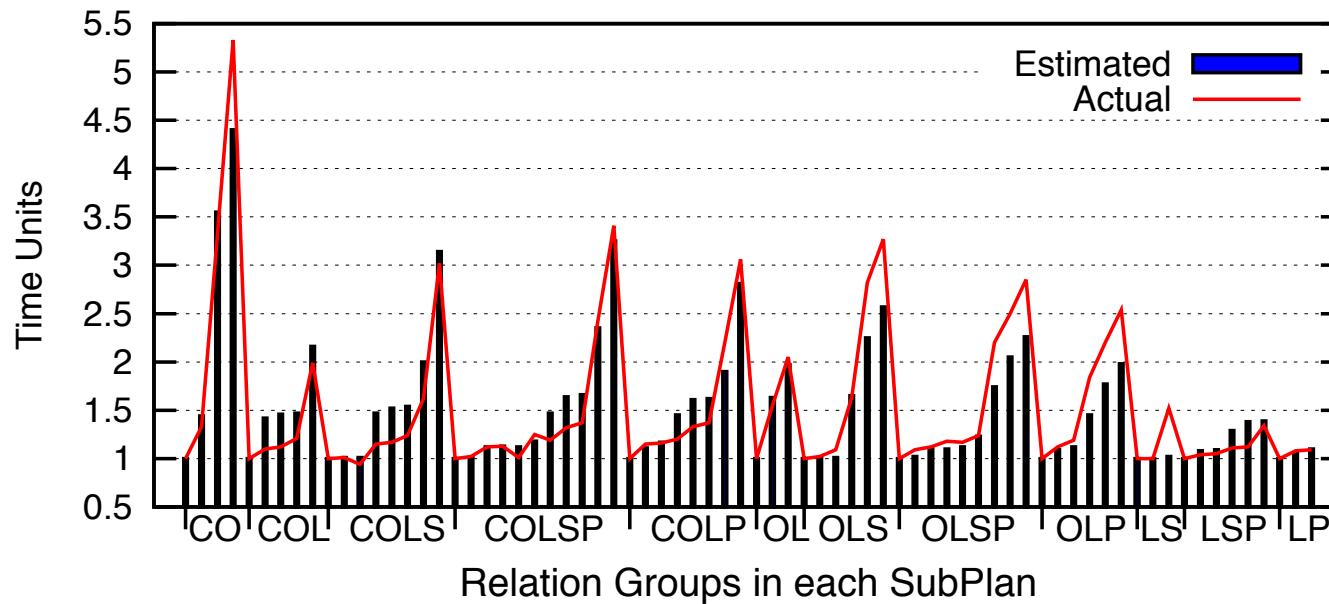
RESULTS: EFFECT OF MULTI-THREADING



Workload: Varying the degree of parallelism

- **Medium-reuse:** 64 queries
- **Thread:** 2-20 threads

RESULTS: ACCURACY OF COST MODEL



Workload:

- **5-way join:** Customer, Orders, Lineitem, Supplier, Part
- **For each sub-plan:** enumerated all HTs that qualify

SUMMARY AND OUTLOOK

Main idea:

- Reuse internal data structures of operators
- More robust towards different workloads

Other topics not covered in this talk

- Reuse-aware hash aggregate
- Top-down plan enumeration algorithm
- Query batch interface

Full paper at **SIGMOD 2017**