

# Co-Evolution of Data and Applications

by

Michael Stonebraker, Dong Deng and  
Michael Brodie

# Our Thesis

- Traditional database design is all wrong
  - Gold standard techniques are not used “in the wild” and for good reasons
  - Contributes to database decay
- If you used the gold standard techniques...
  - You would get application decay
- Probably better to take your poison in both dimensions

# Traditional Database Design Wisdom

- Use a modeling tool (e.g. an Entity-Relationship diagram tool)
- To construct/modify an E-R diagram
- When satisfied, push a button
- Tool spits out a relational schema in 3<sup>rd</sup> normal form
- Code in ODBC/JDBC against this schema

# Example



Two entities (Supplier and Part with attributes)

One relationship (Supply with attributes, which is 1 – N)

Code in ODBC/JDBC for the resulting tables

## Now Suppose....

- Management decides to allow multiple suppliers for each part, as long as they are in different regions
  - Perhaps to get better terms
  - Or ...
- Such changes happen frequently “in the wild”
  - Often once a quarter or more

# New Diagram



Two entities (Supplier and Part with attributes)

One relationship (Supply with attributes which is now M – N)

New tables to code against

# Summary of The Traditional Wisdom

- Convert “old” database to “new”
- Define “old” as views
- Applications (in theory) continue to run
- Database stays in 3NF (defined as goodness by research community)

# A Dirty Little Secret

- Based on a survey of about 20 DBAs at 3 large companies.....
  - None use this methodology  
or
  - Use it for “green field” design and then abandon it
- I.e. the “gold standard” is not used “in the wild”




# Why Not????

- View update problems
- Semantics change – often requires recoding
  - As an aside: few DBAs make any “optional” changes
  - And most research deals with “non semantic” changes
- Net result -- Substantial risk!
  - Applications all over the enterprise must be found and corrected
  - Often with no budget for this activity

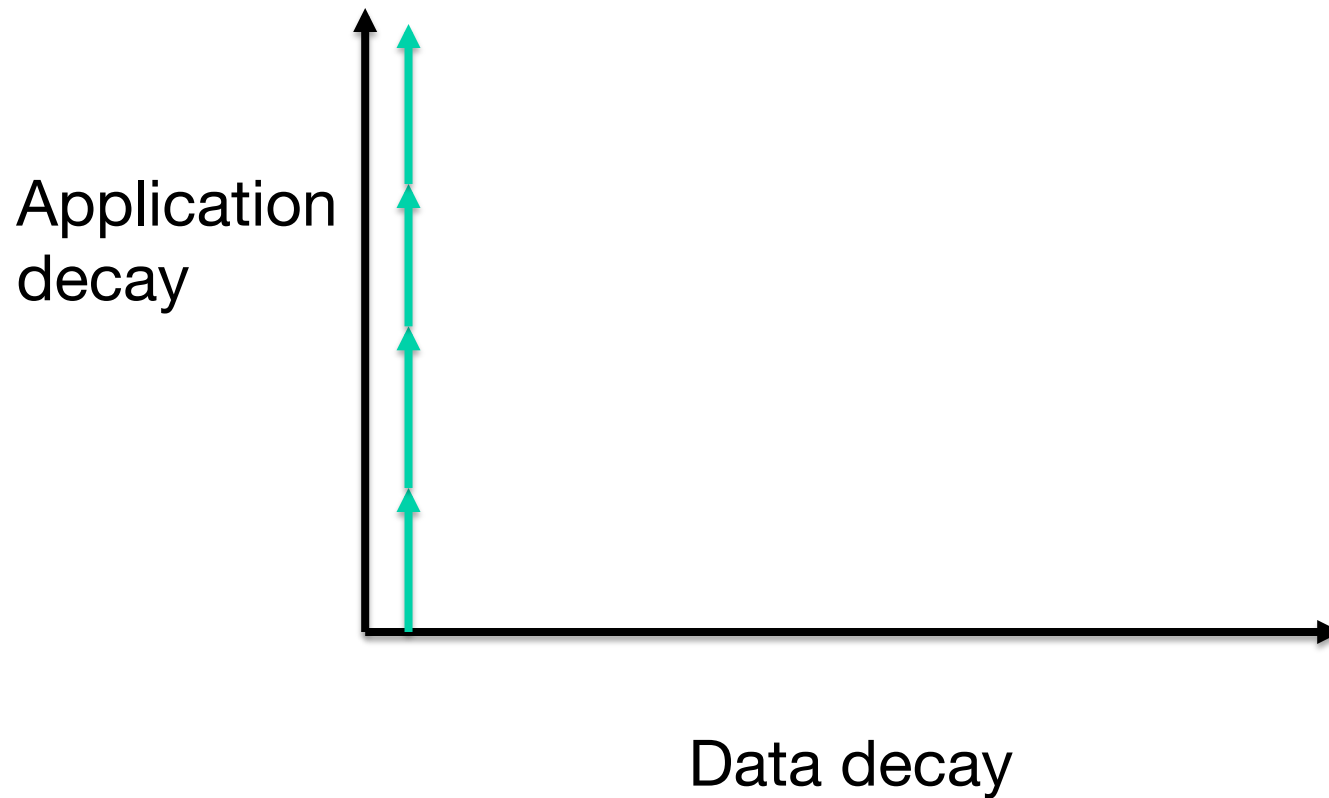
# Less Risky Solution -- Kluge

- Leave the old schema as is
  - And kluge the semantics
  - Often lowers or removes application maintenance!!!!
- But the result does not conform to any ER diagram!
  - No longer 3NF
- Database decays over time!!!!
  - i.e drifts further and further from 3NF


# Looked at Graphically Over 4 Evolutions



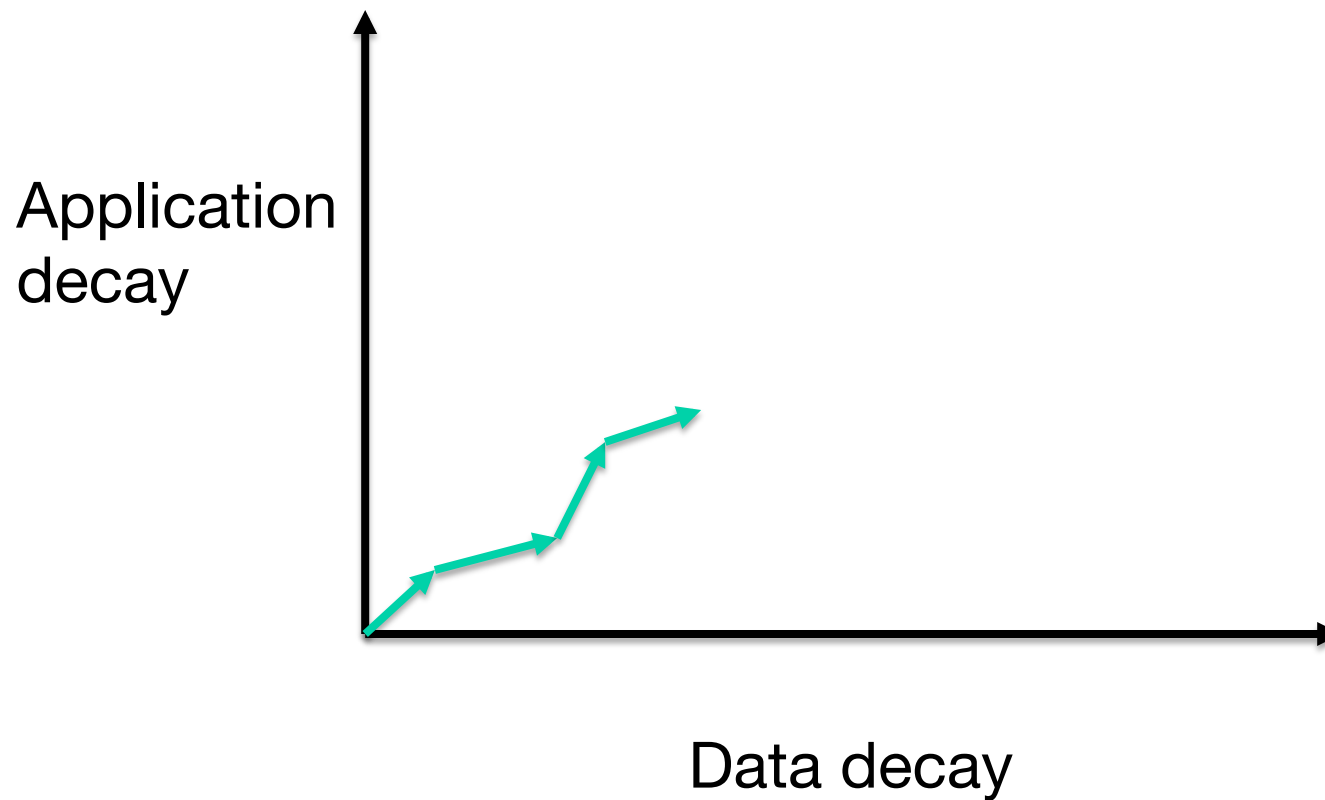
# Suppose You Followed the DBMS Gold Standard....



# Choices So Far.....



# Our Thesis - Co-Evolution Likely to Result in a Better Solution



# We Have....

- An architecture to support co-evolution
  - Cannot allow application groups to independently code in ODBC/JDBC
- Composite cost function for overall decay
  - I.e. evolution cost
- Algorithms to monitor database decay
- Working on algorithms to monitor application decay
- Working on algorithms to suggest schema changes

# HELP!!!!!!!

- We need a real-world DBMS application to try our stuff on
- We need:
  - Your schema and application code
  - Over a few changes